Table 1: Base Models

model	model_parameters	FS Percentile	F1 Train Score	F1 Validation Score
MLPClassifier	hidden_layer_sizes=(20, 20), learning_rate_init=0.01, max_iter=1000	90	0.39 ± 0.02	0.37 ± 0.01
MLPClassifier	hidden_layer_sizes=(20, 20), learning_rate_init=0.01, max_iter=1000	70	0.37 ± 0.02	0.36 ± 0.02
MLPClassifier	hidden_layer_sizes=(20, 20), learning_rate_init=0.01, max_iter=1000	50	0.36 ± 0.02	0.34 ± 0.01
LogisticRegression	class_weight='balanced', max_iter=1000	90	0.32 ± 0.0	0.3 ± 0.01
LogisticRegression	class_weight='balanced', max_iter=1000	70	0.31 ± 0.01	0.3 ± 0.01
LogisticRegression	class_weight='balanced', max_iter=1000	50	0.3 ± 0.01	0.29 ± 0.02
KNeighborsClassifier	n_neighbors=35, p=1, weights='distance'	50	1.0 ± 0.0	0.34 ± 0.01
KNeighborsClassifier	n_neighbors=35, p=1, weights='distance'	70	1.0 ± 0.0	0.34 ± 0.01
KNeighborsClassifier	n_neighbors=35, p=1, weights='distance'	90	1.0 ± 0.0	0.34 ± 0.0
GaussianNB	base model	50	0.27 ± 0.01	0.27 ± 0.01
GaussianNB	base model	70	0.18 ± 0.04	0.18 ± 0.03
GaussianNB	base model	90	0.13 ± 0.01	0.12 ± 0.01
DecisionTreeClassifier	class_weight='balanced', max_depth=15	90	0.46 ± 0.01	0.34 ± 0.03
DecisionTreeClassifier	class_weight='balanced', max_depth=15	70	0.46 ± 0.01	0.34 ± 0.03
DecisionTreeClassifier	class_weight='balanced', max_depth=15	50	0.45 ± 0.01	0.33 ± 0.04

Table 2: Ensembles

model	model_parameters	F1 Train Score	F1 Validation Score
RandomForestClassifier	max_depth: 10, min_samples_split: 2, n_estimators: 100	0.36 ± 0.01	0.35 ± 0.0
RandomForestClassifier	max_depth: 10, min_samples_split: 2, n_estimators: 200	0.36 ± 0.01	0.35 ± 0.0
RandomForestClassifier	max_depth: 10, min_samples_split: 5, n_estimators: 100	0.36 ± 0.0	0.35 ± 0.0
RandomForestClassifier	max_depth: 10, min_samples_split: 5, n_estimators: 200	0.36 ± 0.01	0.35 ± 0.0
RandomForestClassifier	max_depth: 5, min_samples_split: 2, n_estimators: 200	0.25 ± 0.01	0.25 ± 0.01
RandomForestClassifier	max_depth: 5, min_samples_split: 5, n_estimators: 100	0.25 ± 0.01	0.25 ± 0.01
RandomForestClassifier	max_depth: 5, min_samples_split: 5, n_estimators: 200	0.25 ± 0.01	0.25 ± 0.01
RandomForestClassifier	max_depth: 5, min_samples_split: 2, n_estimators: 100	0.25 ± 0.01	0.25 ± 0.0
GradientBoostingClassifier	learning_rate: 0.05, max_depth: 5, n_estimators: 100	0.52 ± 0.03	0.4 ± 0.01
GradientBoostingClassifier	learning_rate: 0.1, max_depth: 5, n_estimators: 50	0.46 ± 0.01	0.4 ± 0.01
GradientBoostingClassifier	learning_rate: 0.1, max_depth: 5, n_estimators: 100	0.48 ± 0.01	0.4 ± 0.01
GradientBoostingClassifier	learning_rate: 0.05, max_depth: 5, n_estimators: 50	0.49 ± 0.02	0.39 ± 0.01
GradientBoostingClassifier	learning_rate: 0.05, max_depth: 10, n_estimators: 50	0.7 ± 0.01	0.39 ± 0.01
GradientBoostingClassifier	learning_rate: 0.05, max_depth: 10, n_estimators: 100	0.75 ± 0.01	0.39 ± 0.01
GradientBoostingClassifier	learning_rate: 0.1, max_depth: 10, n_estimators: 50	0.72 ± 0.01	0.39 ± 0.01
GradientBoostingClassifier	learning_rate: 0.1, max_depth: 10, n_estimators: 100	0.8 ± 0.01	0.38 ± 0.02
BaggingClassifier	estimator=MLPClassifier(hidden_layer_sizes=(20, 20), learning_rate_init=0.01 max_iter=1000, max_features: 1.0, max_samples: 0.8, n_estimators=10	0.4 ± 0.01	0.38 ± 0.02
BaggingClassifier	estimator=MLPClassifier(hidden_layer_sizes=(20, 20), learning_rate_init=0.01 max_iter=1000, max_features: 0.8, max_samples: 0.8, n_estimators=10	0.38 ± 0.01	0.37 ± 0.01
BaggingClassifier	estimator=MLPClassifier(hidden_layer_sizes=(20, 20), learning_rate_init=0.01 max_iter=1000, max_features: 1.0, max_samples: 1.0, n_estimators=10	0.39 ± 0.02	0.37 ± 0.0
BaggingClassifier	estimator=MLPClassifier(hidden_layer_sizes=(20, 20), learning_rate_init=0.01 max_iter=1000, max_features: 0.8, max_samples: 1.0, n_estimators=10	0.36 ± 0.01	0.36 ± 0.01
AdaBoostClassifier	learning_rate: 0.05, n_estimators: 50	0.2 ± 0.0	0.2 ± 0.0
AdaBoostClassifier	learning_rate: 0.05, n_estimators: 100	0.2 ± 0.0	0.2 ± 0.0
AdaBoostClassifier	learning_rate: 0.1, n_estimators: 50	0.2 ± 0.0	0.2 ± 0.0
AdaBoostClassifier	learning_rate: 0.1, n_estimators: 100	0.2 ± 0.0	0.2 ± 0.0

Table 3: MLP Optimization (Top 15)

Activation Function	Architecture	Learning Rate	Initial Learning Rate	Solver	F1 Train Score	F1 Validation Score
relu	(30, 30)	adaptive	0.02	sgd	0.42 ± 0.01	0.4 ± 0.02
relu	(30, 30)	constant	0.02	sgd	0.42 ± 0.01	0.4 ± 0.01
tanh	(30, 30)	constant	0.01	sgd	0.41 ± 0.01	0.4 ± 0.01
relu	(20, 20, 20)	adaptive	0.02	sgd	0.41 ± 0.01	0.39 ± 0.02
relu	(30, 30)	adaptive	0.01	adam	0.41 ± 0.01	0.39 ± 0.01
tanh	(20, 20, 20)	adaptive	0.00	sgd	0.41 ± 0.01	0.39 ± 0.01
relu	(30, 30)	adaptive	0.01	sgd	0.41 ± 0.01	0.39 ± 0.02
relu	(30, 30)	adaptive	0.00	sgd	0.41 ± 0.01	0.39 ± 0.01
relu	(30, 30)	constant	0.00	sgd	0.41 ± 0.01	0.39 ± 0.01
tanh	(20, 20, 20)	constant	0.01	sgd	0.41 ± 0.01	0.39 ± 0.02
tanh	(20, 20, 20)	constant	0.02	sgd	0.41 ± 0.01	0.39 ± 0.02
relu	(30, 30)	constant	0.01	sgd	0.41 ± 0.01	0.39 ± 0.01
logistic	(30, 30)	adaptive	0.00	adam	0.41 ± 0.01	0.39 ± 0.01
tanh	(30, 30)	adaptive	0.00	sgd	0.41 ± 0.01	0.39 ± 0.02
tanh	(30, 30)	adaptive	0.02	sgd	0.42 ± 0.0	0.39 ± 0.01

Table 4: Ensembles Optimization - Bagging

model	model_parameters	F1 Train Score	F1 Validation Score
BaggingClassifier	estimator=MLPClassifier(activation=tanh, hidden_layer_sizes=(30, 30), learning_rate_init=0.01 max_iter=1000, solver=sgd , max_features: 1.0, max_samples: 0.2, n_estimators: 10, n_estimators=10	0.4 ± 0.01	0.38 ± 0.01
BaggingClassifier	estimator=MLPClassifier(activation=tanh, hidden_layer_sizes=(30, 30), learning_rate_init=0.01 max_iter=1000, solver=sgd , max_features: 1.0, max_samples: 0.4, n_estimators: 30, n_estimators=10	0.4 ± 0.01	0.38 ± 0.01
BaggingClassifier	estimator=MLPClassifier(activation=tanh, hidden_layer_sizes=(30, 30), learning_rate_init=0.01 max_iter=1000, solver=sgd , max_features: 1.0, max_samples: 0.4, n_estimators: 50, n_estimators=10	0.41 ± 0.01	0.38 ± 0.01
BaggingClassifier	estimator=MLPClassifier(activation=tanh, hidden_layer_sizes=(30, 30), learning_rate_init=0.01 max_iter=1000, solver=sgd , max_features: 0.8, max_samples: 0.4, n_estimators: 10, n_estimators=10	0.39 ± 0.0	0.37 ± 0.01
BaggingClassifier	estimator=MLPClassifier(activation=tanh, hidden_layer_sizes=(30, 30), learning_rate_init=0.01 max_iter=1000, solver=sgd , max_features: 0.8, max_samples: 0.4, n_estimators: 30, n_estimators=10	0.39 ± 0.01	0.37 ± 0.01
BaggingClassifier	estimator=MLPClassifier(activation=tanh, hidden_layer_sizes=(30, 30), learning_rate_init=0.01 max_iter=1000, solver=sgd , max_features: 1.0, max_samples: 0.4, n_estimators: 10, n_estimators=10	0.4 ± 0.01	0.37 ± 0.01
BaggingClassifier	estimator=MLPClassifier(activation=tanh, hidden_layer_sizes=(30, 30), learning_rate_init=0.01 max_iter=1000, solver=sgd , max_features: 0.8, max_samples: 0.4, n_estimators: 50, n_estimators=10	0.38 ± 0.02	0.37 ± 0.0
BaggingClassifier	estimator=MLPClassifier(activation=tanh, hidden_layer_sizes=(30, 30), learning_rate_init=0.01 max_iter=1000, solver=sgd , max_features: 1.0, max_samples: 0.2, n_estimators: 30, n_estimators=10	0.39 ± 0.02	0.37 ± 0.0
BaggingClassifier	estimator=MLPClassifier(activation=tanh, hidden_layer_sizes=(30, 30), learning_rate_init=0.01 max_iter=1000, solver=sgd , max_features: 1.0, max_samples: 0.2, n_estimators: 50, n_estimators=10	0.39 ± 0.01	0.37 ± 0.0
BaggingClassifier	estimator=MLPClassifier(activation=tanh, hidden_layer_sizes=(30, 30), learning_rate_init=0.01 max_iter=1000, solver=sgd , max_features: 0.8, max_samples: 0.2, n_estimators: 30, n_estimators=10	0.37 ± 0.01	0.36 ± 0.0
BaggingClassifier	estimator=MLPClassifier(activation=tanh, hidden_layer_sizes=(30, 30), learning_rate_init=0.01 max_iter=1000, solver=sgd , max_features: 0.8, max_samples: 0.2, n_estimators: 50, n_estimators=10	0.37 ± 0.01	0.36 ± 0.0
BaggingClassifier	estimator=MLPClassifier(activation=tanh, hidden_layer_sizes=(30, 30), learning_rate_init=0.01 max_iter=1000, solver=sgd , max_features: 0.8, max_samples: 0.2, n_estimators: 10, n_estimators=10	0.36 ± 0.0	0.35 ± 0.0

Table 5: Ensembles Optimization - GradientBoost

model	model_parameters	F1 Train Score	F1 Validation Score
GradientBoostingClassifier	learning_rate: 0.1, max_depth: 5, max_features: 0.5, n_estimators: 50, subsample: 1	0.47 ± 0.01	0.39 ± 0.01
GradientBoostingClassifier	learning_rate: 0.1, max_depth: 5, max_features: 0.5, n_estimators: 100, subsample: 1	0.45 ± 0.05	0.38 ± 0.03
GradientBoostingClassifier	learning_rate: 0.1, max_depth: 5, max_features: 0.5, n_estimators: 50, subsample: 0.5	0.4 ± 0.03	0.38 ± 0.01
GradientBoostingClassifier	learning_rate: 0.1, max_depth: 10, max_features: 0.5, n_estimators: 50, subsample: 1	0.74 ± 0.01	0.37 ± 0.03
GradientBoostingClassifier	learning_rate: 0.1, max_depth: 10, max_features: 0.5, n_estimators: 100, subsample: 1	0.77 ± 0.03	0.37 ± 0.02
GradientBoostingClassifier	learning_rate: 0.1, max_depth: 5, max_features: 0.5, n_estimators: 100, subsample: 0.5	0.43 ± 0.01	0.37 ± 0.01
GradientBoostingClassifier	learning_rate: 0.1, max_depth: 10, max_features: 0.5, n_estimators: 50, subsample: 0.5	0.53 ± 0.0	0.37 ± 0.01
GradientBoostingClassifier	learning_rate: 0.1, max_depth: 5, max_features: 1, n_estimators: 100, subsample: 1	0.43 ± 0.04	0.36 ± 0.01
GradientBoostingClassifier	learning_rate: 0.1, max_depth: 10, max_features: 1, n_estimators: 50, subsample: 1	0.59 ± 0.05	0.36 ± 0.01
GradientBoostingClassifier	learning_rate: 0.1, max_depth: 10, max_features: 1, n_estimators: 100, subsample: 1	0.55 ± 0.12	0.36 ± 0.01
GradientBoostingClassifier	learning_rate: 0.1, max_depth: 10, max_features: 1, n_estimators: 100, subsample: 0.5	0.45 ± 0.0	0.34 ± 0.02
GradientBoostingClassifier	learning_rate: 0.1, max_depth: 10, max_features: 1, n_estimators: 50, subsample: 0.5	0.42 ± 0.01	0.34 ± 0.01
GradientBoostingClassifier	learning_rate: 0.1, max_depth: 5, max_features: 1, n_estimators: 50, subsample: 1	0.37 ± 0.02	0.33 ± 0.0
GradientBoostingClassifier	learning_rate: 0.1, max_depth: 10, max_features: 0.5, n_estimators: 100, subsample: 0.5	0.52 ± 0.05	0.31 ± 0.06
GradientBoostingClassifier	learning_rate: 0.5, max_depth: 5, max_features: 1, n_estimators: 50, subsample: 1	0.34 ± 0.03	0.31 ± 0.04
GradientBoostingClassifier	learning_rate: 0.1, max_depth: 5, max_features: 1, n_estimators: 50, subsample: 0.5	0.35 ± 0.02	0.31 ± 0.01
GradientBoostingClassifier	learning_rate: 0.5, max_depth: 10, max_features: 1, n_estimators: 50, subsample: 1	0.42 ± 0.03	0.29 ± 0.03
GradientBoostingClassifier	learning_rate: 0.5, max_depth: 5, max_features: 0.5, n_estimators: 50, subsample: 1	0.28 ± 0.08	0.27 ± 0.07
GradientBoostingClassifier	learning_rate: 0.5, max_depth: 10, max_features: 1, n_estimators: 50, subsample: 0.5	0.33 ± 0.01	0.26 ± 0.03
GradientBoostingClassifier	learning_rate: 0.1, max_depth: 5, max_features: 1, n_estimators: 100, subsample: 0.5	0.27 ± 0.15	0.25 ± 0.15
GradientBoostingClassifier	learning_rate: 0.5, max_depth: 10, max_features: 1, n_estimators: 100, subsample: 1	0.33 ± 0.01	0.25 ± 0.03
GradientBoostingClassifier	learning_rate: 0.5, max_depth: 5, max_features: 1, n_estimators: 100, subsample: 1	0.27 ± 0.06	0.24 ± 0.04
GradientBoostingClassifier	learning_rate: 0.5, max_depth: 5, max_features: 0.5, n_estimators: 50, subsample: 0.5	0.22 ± 0.12	0.21 ± 0.11
GradientBoostingClassifier	learning_rate: 0.5, max_depth: 10, max_features: 1, n_estimators: 100, subsample: 0.5	0.26 ± 0.05	0.2 ± 0.02
GradientBoostingClassifier	learning_rate: 0.5, max_depth: 5, max_features: 1, n_estimators: 50, subsample: 0.5	0.22 ± 0.06	0.18 ± 0.05
GradientBoostingClassifier	learning_rate: 0.5, max_depth: 10, max_features: 0.5, n_estimators: 50, subsample: 1	0.21 ± 0.04	0.18 ± 0.01
GradientBoostingClassifier	learning_rate: 0.5, max_depth: 10, max_features: 0.5, n_estimators: 50, subsample: 0.5	0.2 ± 0.12	0.16 ± 0.06
GradientBoostingClassifier	learning_rate: 0.5, max_depth: 5, max_features: 1, n_estimators: 100, subsample: 0.5	0.19 ± 0.08	0.16 ± 0.04
GradientBoostingClassifier	learning_rate: 0.5, max_depth: 10, max_features: 0.5, n_estimators: 100, subsample: 0.5	0.17 ± 0.03	0.15 ± 0.04
GradientBoostingClassifier	learning_rate: 0.5, max_depth: 5, max_features: 0.5, n_estimators: 100, subsample: 1	0.14 ± 0.11	0.14 ± 0.1
GradientBoostingClassifier	learning_rate: 0.5, max_depth: 5, max_features: 0.5, n_estimators: 100, subsample: 0.5	0.13 ± 0.03	0.12 ± 0.05
GradientBoostingClassifier	learning_rate: 0.5, max_depth: 10, max_features: 0.5, n_estimators: 100, subsample: 1	0.14 ± 0.03	0.12 ± 0.01